Yajat Mehta

+1 (623) 273 4531 | official.yajatmehta@gmail.com | LinkedIn | Tempe, AZ

Education

Arizona State University — August 2022 – May 2026

B.S., Business Data Analytics | B.S., Artificial Intelligence in Business | GPA: 3.56 — *Tempe, AZ*

Skills

Intelligence & Analysis: NLP, Text Classification, Sentiment Analysis, Entity Extraction, Pattern Recognition, Anomaly Detection, Open-Source Data Analysis

AI & ML: Python, Scikit-learn, TensorFlow, PyTorch, PySpark, Word2Vec, Random Forest, Logistic Regression, NLP Pipelines, Spark ML

Cloud & Data Infrastructure: Google Cloud Platform, Google Cloud Storage, PostgreSQL, Redis, Docker, REST APIs, FastAPI

Security & Governance: Authentication Architecture, OAuth 2.0, Role-Based Access Control, COBIT, NIST CSF, ISO 27001, BCP/DRP, Risk Mitigation, FERPA

Analytics & BI: SQL, Tableau, Power BI, Excel, Data Visualization, Geospatial Data Analysis, Dashboarding

Tools: Git, Postman, VS Code, Google Colab, Figma

Work Experience

Founder & Developer — Rumore, Campus Social Platform — 2025 – Present

FastAPI, Supabase, PostgreSQL, Redis, React Native/Expo, Google OAuth — *Tempe, AZ*

- Designed and implemented a distributed backend system with **Google OAuth 2.0**, **Redis**-based token management, campus-scoped access control, and strict identity separation across user populations — demonstrating systems design under real security constraints.
- Built protected API infrastructure with session lifecycle controls, moderation pipelines, and role-enforced data boundaries — directly transferable to multi-tenant secure application environments.
- Applied privacy-by-design principles to handle sensitive user data, separating anonymous and identified identity layers to minimize exposure and enforce information access controls.

Projects

Sentiment Analysis Using GCP and PySpark — NLP Pipeline — May 2025

- Built a scalable **PySpark** NLP pipeline on **Google Cloud Platform** to classify sentiment across a large unstructured text corpus — applying tokenization, stopword removal, **Word2Vec** vectorization, and **Random Forest** classification.
- Designed for reusability: persisted the trained Spark ML model to **GCS** and evaluated pipeline performance across preprocessing variants — demonstrating production-oriented text intelligence system design.

IT Governance & Risk Mitigation — Jefford's InfoSec Framework — Spring 2026

- Led a team of six to design an 8-quarter risk mitigation roadmap verified against **COBIT**, **ISO 27001**, and **NIST CSF** standards — covering BCP/DRP, endpoint encryption, automated patching, and fraud protection across EU, U.S., and international regulatory jurisdictions.
- Produced phased Gantt-based implementation timelines mapping organizational risk tolerance to initiative sequencing — directly applicable to program management in defense acquisition and compliance environments.

Marketing Assistant AI Agent — Autonomous Decision Support System — May 2026

- Designed a **Python**-based AI agent architecture for multi-source data aggregation, automated insight generation, and structured report production — with explicit human-in-the-loop controls at every decision stage.
- Scoped system failure modes including hallucinated outputs, incomplete retrieval, and accountability gaps — and designed mitigations into the architecture prior to deployment, modeling responsible AI system design under operational constraints.

Factors Influencing Heart Disease — Predictive Classification — December 2024

- Trained **Logistic Regression** and **Random Forest** classifiers using **GridSearchCV** hyperparameter tuning; achieved **82% accuracy**, **90% recall**, and **0.86 AUC** — demonstrating applied ML model development and evaluation methodology transferable to predictive risk classification tasks.